
Infinitely-Armed Bandit Algorithms

Matthew Faulkner

Jon Krause

Daniel Rosenberg

Abstract

We have been considering a variation on the standard K -armed bandit problem where the number of arms is infinite. Here we present our progress and plans for the future.

1 Background

In many situations, the number of arms on a K -armed bandit is actually very large. For example, in the case of choosing an ordering of N ads, the number of possible arms to pull is actually $N!$, which is far too large for any standard finite-armed bandit algorithm to handle. Generalizing, it may be the case that the number of arms on a bandit is actually infinite. Examples of this include optimizing parameters of some algorithm or classifier and maximizing a noisy function over a domain containing an infinite number of points, such as a unit hypercube.

The algorithms designed to handle an infinite-armed bandit can be partitioned into those algorithms which merely discretize the domain into a finite number of points, and those which actually take into consideration every point. One natural question to ask is: are these infinite-armed bandit algorithms worth it, or is it more advantageous to stick with a discretization of the domain into a finite number of points and run a standard finite-armed bandit algorithm on them? With that in mind, we implemented the following algorithms: discretized UCB1, discretized ϵ -greedy, discretized Exp3, Zooming[2], and Hierarchical Optimistic Optimization (HOO)[1]. The latter two of these algorithms are designed to work on an infinite-armed bandit, whereas the former three all pick a finite number of points in the domain to run on.

2 Algorithms

2.1 General Framework

The first order of business in implementing any of these algorithms was abstracting everything out so that they could all be easily applied without too much redundant code. For example, the HOO algorithm works on reward functions over topological spaces, and the Zooming algorithm works on reward functions over metric spaces, so ideally we should just be able to pass in some domain that the problem is over and the algorithm should just be able to work. This is what we have done. In addition, bandits have been abstracted (see section on artificial data).

2.2 Discretization Algorithms

The simplest approach that one can take when tackling an infinite number of arms is to merely pick out some finite number of arms and run a standard finite-armed bandit algorithm on them. This raises the following questions:

- How many arms should be chosen?
- How should the arms be chosen?

The answer to the first question is algorithm- and problem-specific. For example, consider the case where the reward function $r(x) = x$ defined on the interval $[0, 1]$. With no noise, the ϵ -greedy strategy will converge to the correct arm immediately after sampling each arm once, and thus it is beneficial to choose a very large number of arms. A UCB1 or Exp3 strategy, though, will take significantly longer to converge, so if convergence time is an issue, perhaps fewer arms should be chosen at the expense of some accuracy. On the other hand, for a needle-in-a-haystack type of reward function, we clearly want to include more arms so that we have a better chance of achieving the optimal value.

Concerning how the arms themselves are chosen, two obvious strategies are to either choose them randomly or to choose them in a grid (at regular intervals in the 1D case). Choosing the arms in a deterministic way means that, for a fixed number of arms, we can construct a reward function that makes the algorithm behave terribly. Even given an arbitrary number of arms, by choosing a reward function with a maximum near an endpoint of an interval in \mathbb{R} , for example, we can ensure that the algorithm would never do too well. For these reasons, we opt to go with a random choice of arms.

In any case, merely discretizing the domain into a finite number of arms introduces a bias, as our effective hypothesis class (each of the arms) probably no longer contains the optimal hypothesis.

2.3 Zooming Algorithm

If you actually want to consider the infinite number of arms, you need to have certain information about the relationship between them. For instance, arms that are *close* should produce similar results. The Zooming Algorithm is defined to work in any example where the arms form a metric space. In each phase, certain arms are chosen to be *active*, and the algorithm chooses which arm to play from these active arms.

The Zooming Algorithm is composed of multiple phases, each of which is composed of $2^{i_{ph}}$ rounds, where i_{ph} is the current phase number. In a given round, you *activate* an arm if that arm is not covered by another arm. Each arm covers a radius defined by $r_t(v) := \sqrt{8 * i_{ph} / (2 + n_t(v))}$ where v is the active arm, and $n_t(v)$ is the number of times a given arm has been chosen at time t . Each time an arm is played, its radius shrinks, and at the beginning of each round, you *activate* arms until you have a complete covering using a *covering oracle*. This oracle can either return an uncovered arm, or state that there is no such arm. After the space is covered, you play the arm with the optimal index, defined as $I_t(v) := \mu_t(v) + 2 * r_t(v)$

A point of concern to us with this algorithm is that it does not seem to remember anything from previous phases when a new phase starts. The only piece of information it maintains is the phase number, which influences the confidence radius and index, which changes the balance between exploration and exploitation. It seems like it may be better to start on a later phase, given knowledge of the specific problem. Additionally, the *covering oracle* becomes very complicated once you move into more advanced spaces.

Additionally, we have run into some issues under more complicated bandit problems. When there are multiple peaks, the algorithm seems to have trouble converging. This may be due to a bug in our code, or some parameter that needs to be tuned in the algorithm.

2.4 Hierarchical Optimistic Optimization

Our final algorithm addresses the case of infinite arms over an arbitrary topological space X . The hierarchical optimistic optimization algorithm (HOO) applies to this case under several assumptions. First, the payoffs are stochastic (non-adversarial), with distributions of bounded support. Second, the payoff function mapping arms in X to rewards is *weakly Lipschitz*, which in short requires the function to be Lipschitz at any maxima, but does not constrain the function at points away from a maxima.

The HOO algorithm is implemented by iteratively refining a *covering tree* over X , and maintaining confidence bounds for each node in the tree. The covering tree is a binary tree where each node represents (covers) a subset of X equal to the union of all subsets of nodes beneath it in the tree. Whenever the algorithm selects an arm, a new leaf containing that arm is added to the tree. The region covered by a new leaf is a subset of the region covered by its parent, so in this way the leaves define an increasing fine cover of X . For each node n_i in the tree, an empirical average μ_i of all

rewards received by arms covered by node n_i is maintained, along with an upper confidence bound for μ_i . These upper confidence bounds are used whenever the algorithm selects an arm: roughly, the algorithm descends the tree by choosing at each node the child with larger upper confidence bound. In this way, HOO resembles the UCB1 algorithm for gaussian processes, in that nearby points have similar confidence bounds, and that the algorithm optimistically selects using these bounds.

So far, the HOO algorithm has been implemented using a covering tree on the unit hypercube for a Euclidean space. This is specifically analyzed in [1]. While this is a suitable space for many experiments, we are also considering implementing a more specifically “topological” space, such as a torus, which would involve a similarly implemented covering tree, but would allow evaluation of the algorithm on a more interesting space.

3 Artificial Data

For now, we have just been testing the algorithms on reward functions from $[0, 1]$ to \mathbb{R} . We currently have support for polynomials of arbitrary degree, binomial functions, composing two arbitrary functions, and adding noise in the form of multiplying the unbiased result by a random number in a given range.

4 Possible Applications

One issue with bandit algorithms is that there are relatively few situations where a pure bandit algorithm is applicable. Rather, it is typically the case that there is some side information (e.g. a search term) that can be used in determining which arm to choose. In that case, this turns into a contextual bandit problem, which the algorithms we have implemented are not suited for. It might be possible to adapt these algorithms for the contextual bandit problem, but that would appear to be a project in and of itself.

5 Conclusion

In conclusion, we did some stuff, and will do more.

6 Future Work

In the next month or so, we intend to add more domains that our algorithms can work over. This should be relatively easy due to the abstraction as described above, although certain aspects might be difficult (i.e. constructing a covering oracle). We would also like to move beyond artificial data. One possibility is applying these algorithms to locations of a set of sensors. While there is real data available for this, the data is naturally only taken at a finite number of locations or configurations, and so we would have to interpolate this data to come up with a reward function defined over all possible positions.

Another possibility is applying these algorithms to parameters in other algorithms. For example, when designing an AI for say, Tic-Tac-Toe, we might have a scoring function on a given board describing how ‘good’ that board is for the current player, and play the move that gives us the highest score. One (not necessarily optimal) way to implement this scoring function would be to assign a weight to each square on the board, and score the board by the weight of each position combined with which player occupies each position. Here, one of our bandit algorithms could be used to determine the optimal set of weights to use. Since these algorithms do not necessarily work against an adversary, though, we would have to have them play against a random player.

Going back to artificial data, we could also try to construct reward functions that cause certain algorithms to do well and others to do poorly. This would be proof that no algorithm is necessarily better than the others in all cases, and would answer our initial “Is it worth it?” question. We could

also attempt to improve the Zooming algorithm so that information is carried over between the different phases, make our own experimental algorithms, or implement more algorithms from more papers. At this point, we have not decided for sure which of these options we would like to pursue, which is also dependent on how much time they take up.

7 Plots

The following plots were based on the reward function $R(x) = \text{noise}(\text{bin}(3x^2 - 3x + 1), .5, 1)$, where $\text{noise}(f(x), a, b)$ multiplies each value received from $f(x)$ by some number chosen uniformly at random between a and b , and $\text{bin}(f(x))$ is 1 with probability $f(x)$ and 0 otherwise. Our implementation of the ϵ -greedy algorithm samples each arm once at first, then proceeds with $\epsilon = \frac{\text{arms}}{\text{rounds}}$. The Exp3 algorithm uses $\gamma = \frac{\text{arms}}{\text{rounds}+1}$. In each discretization case, 100 arms were used.




Figure 1: ϵ greedy




Figure 2: EXP3




Figure 3: Discretized UCB1

All of our code is available at:

<http://code.google.com/p/bandits/>




Figure 4: HOO algorithm regret.

References

- [1] S. Bubeck, R. Munos, G. Stoltz, and C. Szepesvari. “Online Optimization in χ -Armed Bandits”. *Advances in Neural Information Processing Systems 21*, 2009, pp. 201-208.
- [2] R. Kleinberg, A. Slivkins, and E. Upfal. “Multi-Armed Bandits in Metric Spaces”. *Proceedings of the 40th ACM Symposium on Theory of Computing*, 2008.